

# Mohanad Mahmoud Sayed

1- How many pods exist on the system?

```
controlplane:~$ k get po
No resources found in default namespace.
controlplane:~$
```

2- How many Nodes exist on the system?

→ There is one worker node which is node01

NAME	STATUS	ROLES	AGE	VERSION
controlplane	Ready	control-plane	29d	v1.32.1
node01	Ready	<none>	29d	v1.32.1

```
controlplane:~$
```

3- Create a new pod with the nginx image.  
Image name: nginx

```
controlplane:~$ kubectl run nginx --image=nginx
pod/nginx created
controlplane:~$ k get pods
NAME    READY   STATUS    RESTARTS   AGE
nginx   1/1     Running   0           23s
controlplane:~$
```

4- Which nodes are these pods placed on?

→ On node: node01

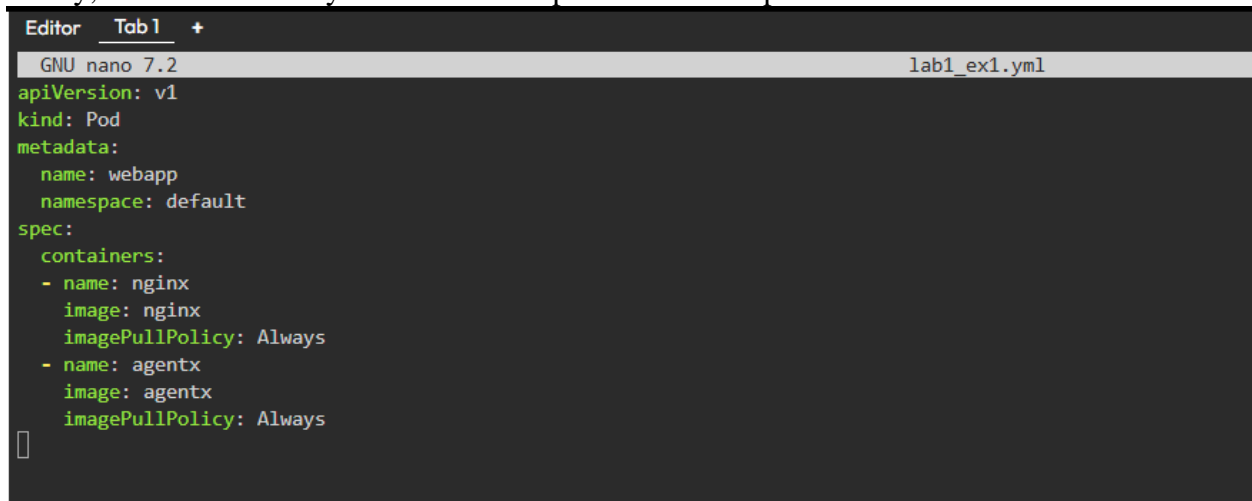
NAME	READY	STATUS	RESTARTS	AGE	IP	NODE	NOMINATED NODE	READINESS GATES
nginx	1/1	Running	0	56s	192.168.1.4	node01	<none>	<none>

```
controlplane:~$
```

5- Create pod from the below yaml using kubectl apply command

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  namespace: default
spec:
  containers:
  - image: nginx
    imagePullPolicy: Always
    name: nginx
  - image: agentx
    imagePullPolicy: Always
    name: agentx
```

Firstly, let's create a file.yml for this example to create the pod

A screenshot of a terminal window showing the nano text editor. The editor is open to a file named 'lab1\_ex1.yml'. The content of the file is a Kubernetes Pod manifest. The top of the editor shows 'Editor Tab 1 +'. The file content is as follows:

```
apiVersion: v1
kind: Pod
metadata:
  name: webapp
  namespace: default
spec:
  containers:
  - name: nginx
    image: nginx
    imagePullPolicy: Always
  - name: agentx
    image: agentx
    imagePullPolicy: Always
```

Then, let's apply

A screenshot of a terminal window showing the execution of kubectl commands. The first command is 'kubectl apply -f lab1\_ex1.yml', which results in 'pod/webapp created'. The second command is 'k get pods', which displays a table of pods. The table has columns for NAME, READY, STATUS, RESTARTS, and AGE. There are two pods listed: 'nginx' with 1/1 ready and 'Running' status, and 'webapp' with 1/2 ready and 'ImagePullBackOff' status.

```
controlplane:~$ kubectl apply -f lab1_ex1.yml
pod/webapp created
controlplane:~$ k get pods
NAME      READY   STATUS              RESTARTS   AGE
nginx     1/1     Running             0           10m
webapp    1/2     ImagePullBackOff    0           29s
controlplane:~$
```

6- How many containers are part of the pod `webapp`

➔ Webapp pod has 2 containers: `nginx` and `agentx`

```
controlplane:~$ kubectl get pod webapp -o jsonpath='{.spec.containers[*].name}'
nginx agentxcontrolplane:~$
controlplane:~$
```

7- What images are used in the new `webapp` pod?

➔ `nginx` and `agentx`

(but `nginx` only is used, `agentx` is not pulled from dockerhub as shown in the ss in question 5)

```
controlplane:~$ kubectl get pod webapp -o jsonpath='{.spec.containers[*].image}'
nginx agentxcontrolplane:~$
controlplane:~$
```

8- What is the state of the container `agentx` in the pod `webapp`

```
agentx:
  Container ID:
  Image:        agentx
  Image ID:
  Port:         <none>
  Host Port:    <none>
  State:        Waiting
    Reason:     ImagePullBackOff
  Ready:        False
  Restart Count: 0
  Environment:  <none>
```

◇ **State:** Waiting

◇ **Reason:** ImagePullBackOff

9- Why do you think the container `agentx` in pod `webapp` is in error?

→**Reason:** Kubernetes is trying to pull the image `agentx` but **fails** because it likely does **not exist on a public registry like Docker Hub**

This causes the container to go into the **ImagePullBackOff** state — which means it tried to pull the image, failed, and now backs off from trying again immediately.

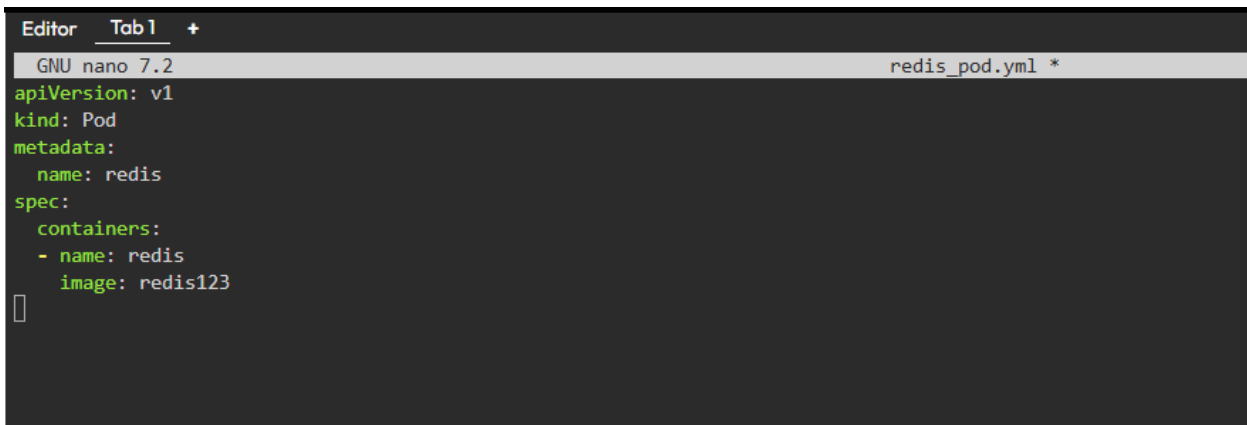
10- Delete the `webapp` Pod.

```
controlplane:~$ kubectl delete pod webapp
pod "webapp" deleted
controlplane:~$
```

11- Create a new pod with the name `redis` and with the image `redis123`.

- Name: `redis`
- Image Name: `redis123`

Firstly, let's create a `file.yml` for this example to create the pod



```
Editor  Tab1  +
GNU nano 7.2                                redis_pod.yml *
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
  - name: redis
    image: redis123

```

Then, let's apply

```
controlplane:~$ kubectl apply -f redis_pod.yml
pod/redis created
controlplane:~$ k get pods
NAME      READY   STATUS             RESTARTS   AGE
nginx     1/1     Running            0           35m
redis     0/1     ContainerCreating  0           3s
controlplane:~$
```

→Note: the image is not running because redis123 is not found on dockerhub to pull

12- Now change the image on this pod to **redis**.

Once done, the pod should be in a **running** state.

Firstly: let's edit the pod file to change the image name from redis123 to redis

```
Editor  Tab1  +
GNU nano 7.2                                redis_pod.yml *
apiVersion: v1
kind: Pod
metadata:
  name: redis
spec:
  containers:
  - name: redis
    image: redis
```

Then, let's apply

```
controlplane:~$ nano redis_pod.yml
controlplane:~$ kubectl apply -f redis_pod.yml
pod/redis configured
controlplane:~$ k get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           38m
redis     1/1     Running   0           2m44s
controlplane:~$
```

✓ The pod is in the running state

13- Create a pod called **my-pod** of image **nginx:alpine**

```
controlplane:~$ kubectl run my-pod --image=nginx:alpine
pod/my-pod created
controlplane:~$ k get pods
NAME      READY   STATUS    RESTARTS   AGE
my-pod    1/1     Running   0           7s
```

14- Delete the pod called my-pod

```
controlplane:~$ k delete pod my-pod
pod "my-pod" deleted
controlplane:~$ k get pods
NAME      READY   STATUS    RESTARTS   AGE
nginx     1/1     Running   0           43m
redis     1/1     Running   0           7m31s
controlplane:~$
```