1- How many `DaemonSets` are created in the cluster in all namespaces?

```
controlplane:~$ kubectl get daemonsets --all-namespaces
NAMESPACE     NAME         DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR            AGE
kube-system   canal        2         2         2       2            2           kubernetes.io/os=linux   35d
kube-system   kube-proxy   2         2         2       2            2           kubernetes.io/os=linux   35d
```

2- what DaemonSets exist on the kube-system namespace?

```
controlplane:~$ kubectl get daemonsets -n kube-system
NAME         DESIRED   CURRENT   READY   UP-TO-DATE   AVAILABLE   NODE SELECTOR            AGE
canal        2         2         2       2            2           kubernetes.io/os=linux   35d
kube-proxy   2         2         2       2            2           kubernetes.io/os=linux   35d
```

3- What is the image used by the POD deployed by the `kube-proxy` DaemonSet

```
      fieldPath: spec.nodeName
image: registry.k8s.io/kube-proxy:v1.32.1
imagePullPolicy: IfNotPresent
```

4- Deploy a DaemonSet for `FluentD` Logging. Use the given specifications.

Name: elasticsearch
Namespace: kube-system
Image: k8s.gcr.io/fluentd-elasticsearch:1.20

```yaml
apiVersion: apps/v1
kind: DaemonSet
metadata:
  name: elasticsearch
  namespace: kube-system
spec:
  selector:
    matchLabels:
      name: elasticsearch
  template:
    metadata:
      labels:
        name: elasticsearch
    spec:
      containers:
      - name: fluentd
        image: k8s.gcr.io/fluentd-elasticsearch:1.20
        resources:
          limits:
            memory: 200Mi
            cpu: 100m
          requests:
            memory: 200Mi
            cpu: 100m
```

```
controlplane:~$ kubectl apply -f fluentd-daemonset.yaml
daemonset.apps/elasticsearch created
```

5- Deploy a pod named `nginx-pod` using the `nginx:alpine` image with the labels set to `tier=backend`.

```
controlplane:~$ kubectl run nginx-pod --image=nginx:alpine --labels=tier=backend --restart=Never
pod/nginx-pod created
```

6- Deploy a `test` pod using the `nginx:alpine` image.

```
controlplane:~$ kubectl run test --image=nginx:alpine --restart=Never
pod/test created
```

7- Create a service `backend-service` to expose the `backend` application within the cluster on port `80`.

```yaml
apiVersion: v1
kind: Service
metadata:
  name: backend-service
spec:
  selector:
    tier: backend
  ports:
  - protocol: TCP
    port: 80
    targetPort: 80
```

```
controlplane:~$ kubectl apply -f backend-service.yaml
service/backend-service created
```

8- try to curl the backend-service from the test pod. What is the response?

```
controlplane:~$ kubectl exec -it test -- /bin/sh
/ # curl backend-service:80
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
/ #
```

9- Create a deployment named `web-app` using the image `nginx` with `2` replicas

```
controlplane:~$ kubectl create deployment web-app --image=nginx --replicas=2
deployment.apps/web-app created
```

10- Expose the `web-app` as service `web-app-service` application on port 80 and nodeport `30082` on the nodes on the cluster

```yaml
apiVersion: v1
kind: Service
metadata:
  name: web-app-service
spec:
  selector:
    app: web-app
  ports:
    - protocol: TCP
      port: 80
      targetPort: 80
      nodePort: 30082
  type: NodePort
```

```
controlplane:~$ kubectl apply -f  web-app-service.yaml
service/web-app-service created
```

11- access the web app from the node

```
controlplane:~$ curl http://172.30.2.2:30082
<!DOCTYPE html>
<html>
<head>
<title>Welcome to nginx!</title>
<style>
html { color-scheme: light dark; }
body { width: 35em; margin: 0 auto;
font-family: Tahoma, Verdana, Arial, sans-serif; }
</style>
</head>
<body>
<h1>Welcome to nginx!</h1>
<p>If you see this page, the nginx web server is successfully installed and
working. Further configuration is required.</p>

<p>For online documentation and support please refer to
<a href="http://nginx.org/">nginx.org</a>.<br/>
Commercial support is available at
<a href="http://nginx.com/">nginx.com</a>.</p>

<p><em>Thank you for using nginx.</em></p>
</body>
</html>
```

12- How many static pods exist in this cluster in all namespaces?

Zero

13-On which nodes are the static pods created currently?

kubectl get pods --all-namespaces -o wide